

Aditya Jhaveri

Jersey City, NJ • adityaajhaveri@gmail.com • (929) 620-9756
adityajhaveri.com • [linkedin.com/in/adityaalokjhaveri](https://www.linkedin.com/in/adityaalokjhaveri) • github.com/swooshie

SUMMARY

Software engineer with **3+ years** of experience building and scaling **distributed systems** in **Java**, **Python**, and **Kafka**, processing 3B+ daily records across production **microservices**. Sole engineer on a production ERP managing 8,000+ devices at NYU. **NYU Tandon M.S. Computer Science** graduate (2026). Open to backend, data engineering, and ML engineering roles.

SKILLS

Languages: Java, Python, JavaScript (ES6+), TypeScript, C++, C#, SQL, Bash, HTML, CSS

Backend & Frameworks: Spring Boot, .NET Core, JPA (Hibernate), REST APIs, MVC, Next.js, Express.js, FastAPI, Swagger, Maven, Node.js, React, Bootstrap

Distributed Systems & Big Data: Apache Kafka, Hadoop (HDFS), Spark, Apache Thrift, Hive, Pig, BigQuery

Databases: PostgreSQL, MySQL (RDBMS), MongoDB, Firestore, Cassandra (NoSQL), Redis

Cloud & DevOps: AWS, Google Cloud (GCP), Docker, GitHub Actions, CI/CD, Git, Version Control, Agile, Vercel

Data, ML & Analytics: TensorFlow, Keras, NumPy, Pandas, SciPy, Scikit-Learn, GeoPandas, OpenCV, QDRANT, UnSloth, LangChain

Tools & Testing: JUnit, PyTest, Postman, JIRA, IntelliJ, PyCharm, Jupyter Notebook, Visual Studio Code, QGIS

WORK EXPERIENCE

Software Engineer | [New York University](#), New York, NY

02/2025 – 05/2026

- **Sole engineer** on NYU GEMSS's production Asset Management System, owning the full codebase, **infrastructure**, **CI/CD pipeline**, and **Dockerized** deployment for a system managing **8,000+ devices** and **300+ employees** across the Admissions department
- Drove phased **AI-assisted migration** from Google Apps Script to a **Node.js/Express**, **TypeScript**, and **MongoDB** architecture, building **MVC-structured REST API** layer with 10+ endpoints and introducing **real-time data sync** while preserving Google Sheets access during transition
- Engineered a **fault-tolerant concurrent** form submission system for employee onboarding/offboarding workflows, cutting retry failures by 40% under peak load; replaced slow service-layer search with a **set-theory-backed** compound filtering engine, improving record lookup speed by 50% across 8,000+ device records

Software Development Engineer | [SAINAPSE](#), India

07/2022 – 05/2024

- Optimized microservice file transfers via **Apache Kafka** by implementing **byte-level serialization/deserialization** and parallel file distribution, reducing time complexity by 50% and space complexity by 33% on **AWS (Linux OS)**
- Solely designed and implemented **HDFS** infrastructure in **Java** to support high-volume ingestion, storing 3B+ daily records with **Hive** across 10+ **microservices**, lifting average performance by 37%
- Designed and shipped a **batch data adapter** attached to **Google BigQuery tables**, enabling on-demand transfer of 2B+ rows and supporting downstream ML workflows such as free-text search and deduplication
- Accelerated data extraction speed from **3GB+ XLSX** and **DOCX** files by updating parsing logic with the Apryse **Java** SDK, improving speed by 30%
- **Mentored** 4 university interns on product architecture and core technology stack, supporting independent feature delivery

Data Science Intern | [AIDASH](#), India

01/2022 – 06/2022

- Architected an end-to-end **LANDSAT** classification pipeline leveraging **ResNet**, **Keras**, and **TensorFlow**; reduced image processing time by 40% while maintaining 95% classification accuracy
- Automated **satellite image labeling** in **QGIS** with a custom **Python** script, reducing manual processing time by 80%

EDUCATION

NEW YORK UNIVERSITY, New York, NY, USA

09/2024 – 05/2026

M.S. in Computer Science | **Coursework:** Algorithms, Machine Learning, Big Data, Search Engines, Options Pricing

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE, Pilani, India

08/2018 – 07/2022

B.E. in Computer Science, Minor in Data Science | **Coursework:** Data Structures, Database Systems, Machine Learning, AI, NLP

PROJECTS

Smart Task Manager ([GitHub](#) | [Live](#))

- Orchestrated a **distributed full-stack system** across 4 services (Next.js, .NET API, MongoDB, Redis, Python) supporting **AI recommendations**, **geolocation-aware reminders**, and real-time cross-service task workflows
- Engineered JWT-secured APIs and **Telegram/Linq webhook** pipelines handling 100+ concurrent operations with sub-200ms average latency

Poker AI - Sealed-State Multi-Agent Simulation ([GitHub](#)) | [NYC Tech Week Hackathon](#), Top 8 Finalist (130 builders, co-hosted by NVIDIA × a16z)

- Architected the **agentic graph** and **backend skeleton** for a sealed-state **multi-agent simulation** where 4 LLM-powered agents compete across 3 scripted Texas Hold hands with fully isolated private state and independent per-agent memory
- Optimized demo reliability by pre-generating all **LLM outputs** (NVIDIA Nemotron 70B) and **Eleven Labs** voice lines offline

Microsoft MARCO Search Engine ([GitHub](#))

- Constructed an 8.8M-passage **search engine** with a compressed inverted index, BM25 scoring, and **DAAT traversal**, achieving 65% storage reduction over a naive index implementation

Portfolio Website & Recruiter Chatbot ([Link](#))

- Launched a **RAG-based recruiter assistant** with a **4-document markdown** knowledge base, bounded query scope to eliminate out-of-context responses, and **round-robin Gemini API** key rotation to handle free-tier rate limits without downtime